

BlinkIt Sales Report

1. Data Sources

Raw Data Files in Original Format

The raw data is stored in **Azure Blob Storage** and mounted to **Databricks** under `/mnt/blinkit_storage/`. The sources include:

- blinkit_customer_feedback.csv
- blinkit_customers.csv
- blinkit_delivery_performance.csv
- blinkit_inventory.csv
- blinkit_marketing_performance.csv
- blinkit_products.csv
- blinkit_orders.csv
- blinkit_order_items.csv

Documentation Describing Sources & Relationships

Each dataset contributes to different business aspects:

- **Customers:** Contains customer profiles and segments.
- **Orders:** Tracks all customer orders.
- **Order Items:** Itemized breakdown of each order.
- **Products:** Product catalog with pricing and categories.
- **Delivery Performance:** Records actual vs. promised delivery times.
- **Marketing Performance:** Details promotional campaigns.
- **Customer Feedback:** Captures ratings and sentiments.
- **Inventory:** Logs stock movements and damages.

Data Dictionary

- `customer_id`: Unique identifier for customers.
- `order_id`: Unique identifier for orders.
- `product_id`: Unique identifier for products.
- `order_total`: Total revenue from an order.
- `delivery_status`: Indicates whether an order was delivered on time.
- `feedback_category`: Category of customer feedback.
- `campaign_name`: Name of marketing campaign.

2. Staging Delta Tables

Staging Tables in Delta Format

The data from **Azure Blob Storage** was loaded into **staging Delta tables** in Databricks:

```
df_staging_customers = spark.read.format("csv") \
    .option("header", "true") \
    .option("inferSchema", "true") \
    .load("dbfs:/mnt/blinkit_storage/blinkit_customers.csv")
df_staging_customers.write.format("delta").mode("overwrite").saveAsTable("stg_customers")
```

Each dataset was stored in **staging Delta tables** for further transformation.

3. ETL Implementation

Complete ETL Code/Workflows

The ETL pipeline includes:

- Reading raw data from **Azure Blob Storage**.
- Storing data in **staging Delta tables**.
- Transforming data for analytical queries.

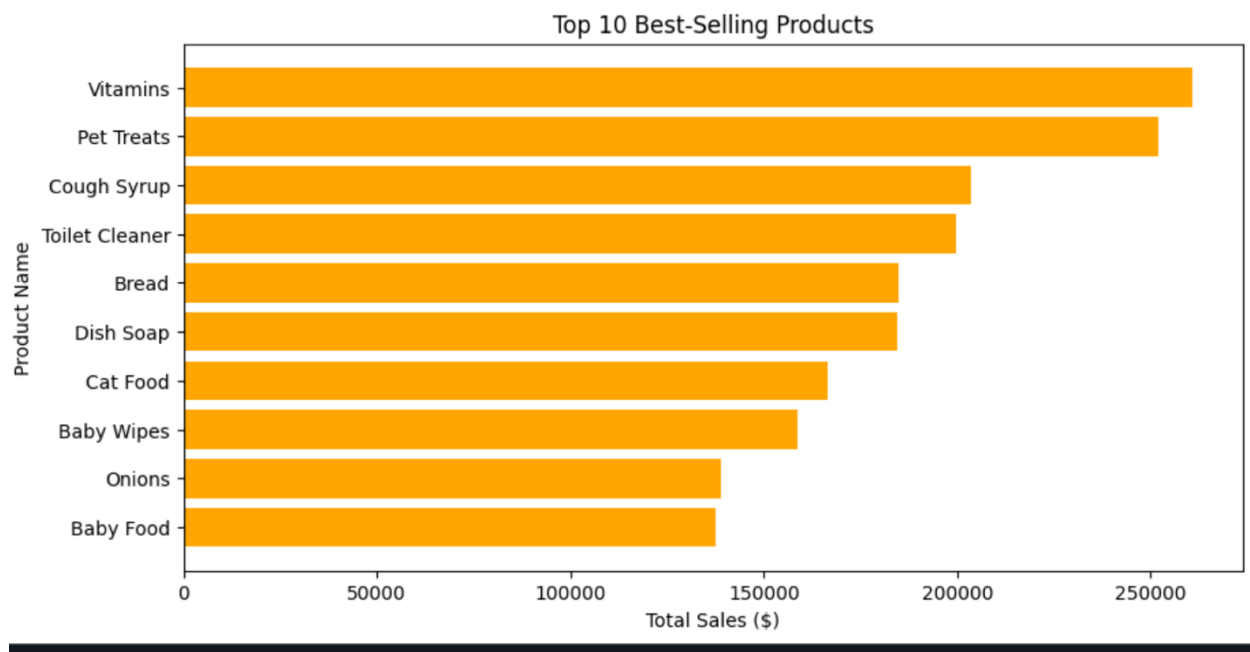
Handling ETL Challenges

- **Data Quality:** Ensured via schema inference and type validation.
- **Slowly Changing Dimensions (SCD Type 2):** Implemented for Customers & Products.
- **Handling Missing Values:** Used fillna() to handle null values.

4. Analytical Queries

Query 1: Top 10 Best Selling Products


```
df_top_products = spark.sql("""
SELECT
    p.product_name,
    SUM(f.quantity) AS total_units_sold,
    SUM(f.quantity * f.unit_price) AS total_sales
FROM fact_orders f
JOIN dim_products p ON f.product_id = p.product_id
GROUP BY p.product_name
ORDER BY total_sales DESC
LIMIT 10
""")
```



Query 2: On-Time vs Delayed Delivery

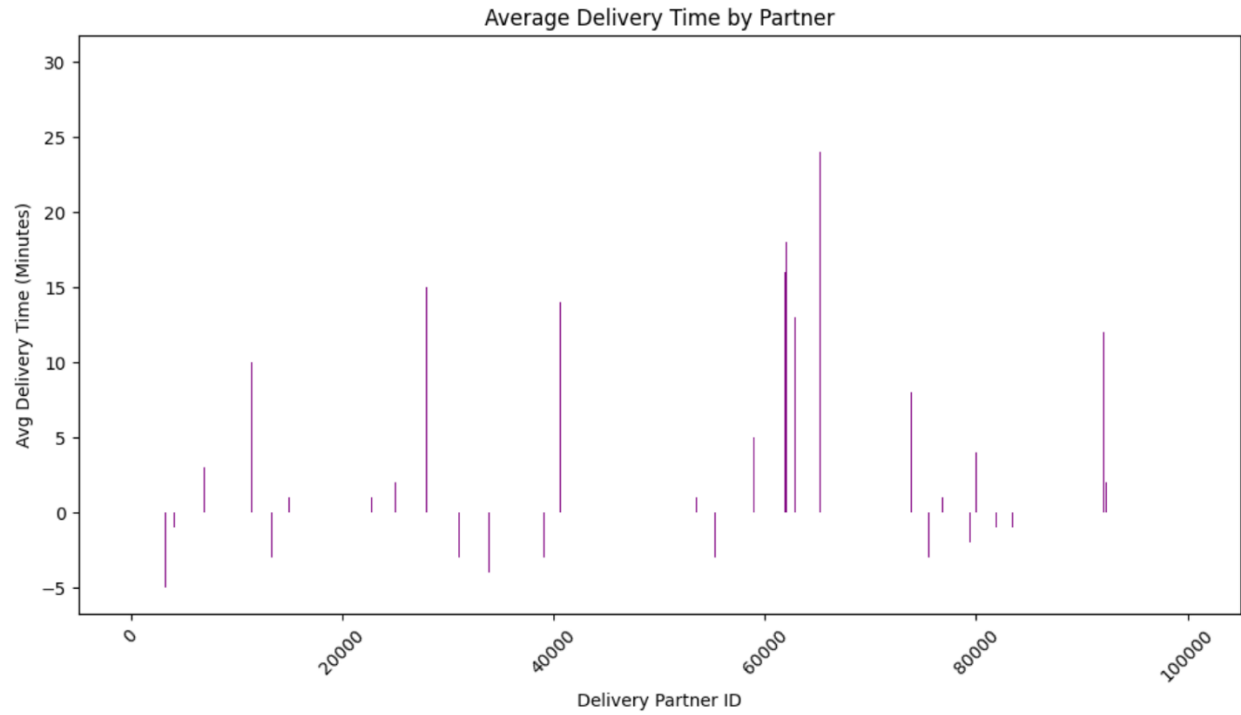
```
df_delivery_performance = spark.sql("""  
SELECT  
    d.delivery_status,  
    COUNT(f.order_id) AS total_orders,  
    AVG(d.delivery_time_minutes) AS avg_delivery_time  
FROM fact_orders f  
JOIN dim_delivery d ON f.delivery_partner_id = d.delivery_partner_id  
GROUP BY d.delivery_status """)
```

Delivery Performance: On-Time vs Delayed



Query 3: Average Delivery Time by Partner


```
df_delivery_time = spark.sql("""  
SELECT  
    d.delivery_partner_id,  
    AVG(d.delivery_time_minutes) AS avg_delivery_time,  
    COUNT(f.order_id) AS total_deliveries  
FROM fact_orders f  
JOIN dim_delivery d ON f.delivery_partner_id = d.delivery_partner_id  
GROUP BY d.delivery_partner_id  
ORDER BY avg_delivery_time  
""")
```



Query 4: Customer Retention

```
df_repeat_customers = spark.sql("""  
    SELECT  
        COUNT(DISTINCT CASE WHEN c.total_orders > 1 THEN c.customer_id END) AS repeat_customers,  
        COUNT(DISTINCT c.customer_id) AS total_customers  
    FROM dim_customers c  
""")
```

Customer Retention: Repeat Orders Analysis



5. Summary of Project Approach & Implementation

This project implemented a **data warehouse** to analyze e-commerce transactions. The pipeline:

1. **Extracted** raw data from **Azure Blob Storage** into Databricks.
2. **Stored** data in **staging Delta tables**.
3. **Transformed** data for analytical queries.
4. **Executed analytical queries** to generate insights.

Challenges & Solutions

- **Handling SCD (Slowly Changing Dimensions):** Implemented SCD Type 2.
- **Ensuring Data Integrity:** Used constraints and data validation.
- **Optimizing Query Performance:** Partitioned fact_orders on time_key.

Effectiveness of the ETL Pipeline

- **Simplified Data Management:** Delta tables provide incremental updates.
- **Performance Optimization:** Queries run efficiently on Databricks.

- **Scalability:** Can accommodate future data growth.

Recommendations for Future Improvements

- **Real-time data streaming** for instant analytics.
 - **Machine learning-based predictions** for customer behavior.
 - **Improved indexing** for faster query performance.
-